Question 3. ***Install Hadoop (anyversion) as a pseudocluster mode and Show Word counting example with your own document file.***

Initialized an instance with Ubuntu 20.04 and downloaded the key.pem of it. Using that file and Putty connected to the server. First thing to do was updating the virtual machine and downloading the necessary softwares. I wanted to use the Hadoop 2.7.7 version and Java 8.0 was compatible with it.

- ***sudo apt install openjdk-8-jdk -y***

and

- ***wget https://archive.apache.org/dist/hadoop/common/hadoop-2.7.7/hadoop-2.7.7.tar.gz***

commands were helpful.

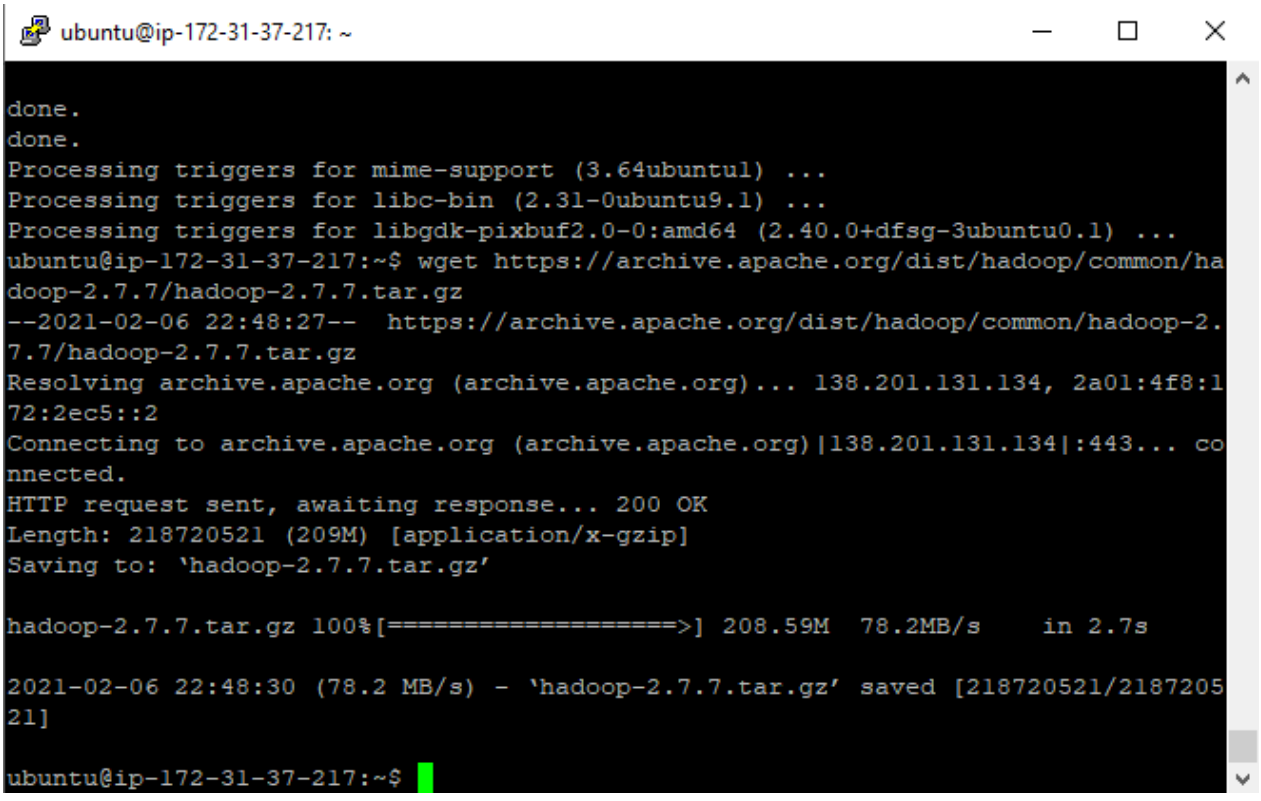

After getting the tar file just untarred it and **deleted** the tar file since dont needed it anymore.

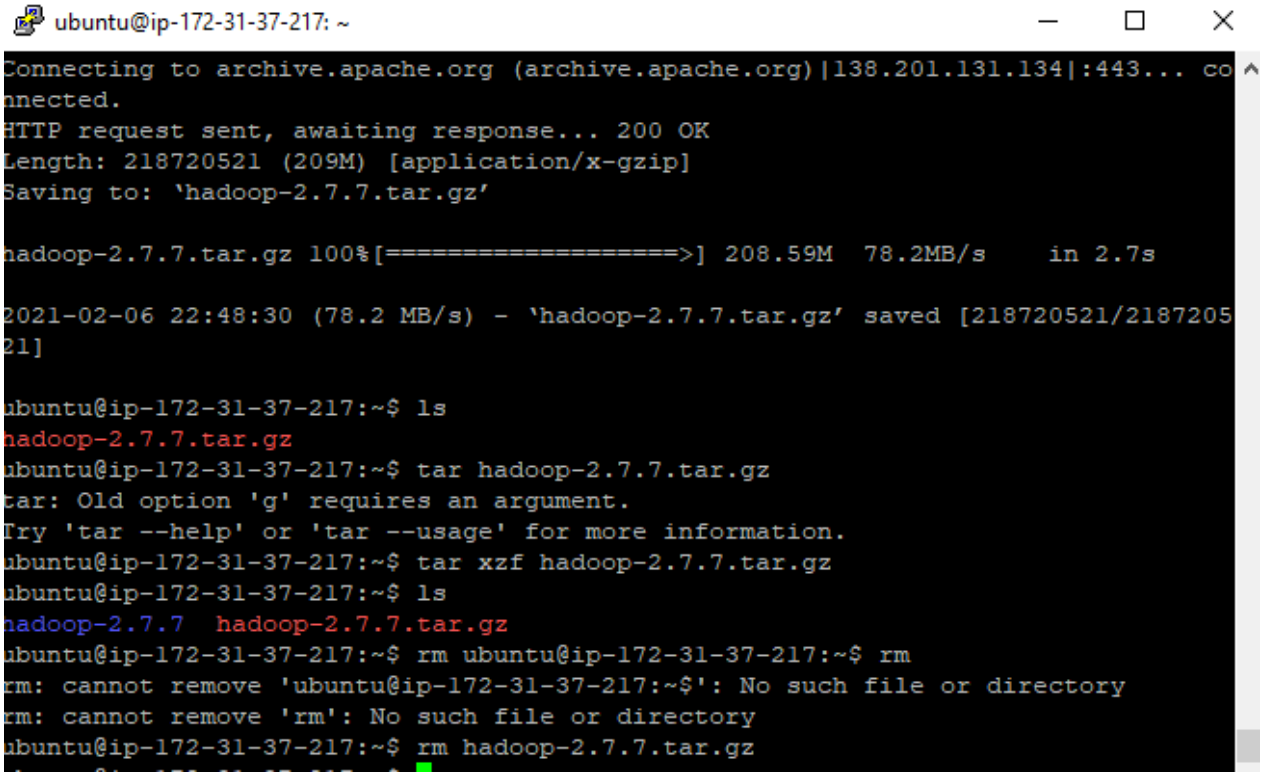

Edited the .**bashrc** shell configuration file to define the Hadoop **environment variables**.

Added these 3 lines to hadoop environment variables to work compatible with java:

- ***export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64***
- ***export PATH=${JAVA_HOME}/bin:${PATH}***
- ***export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar***

To set up Hadoop in a pseudo-distributed mode, specified the URL for NameNode. Opened the *core-site.xml* file.

- ***sudo nano $HADOOP_HOME /etc/hadoop/core-site.xml***

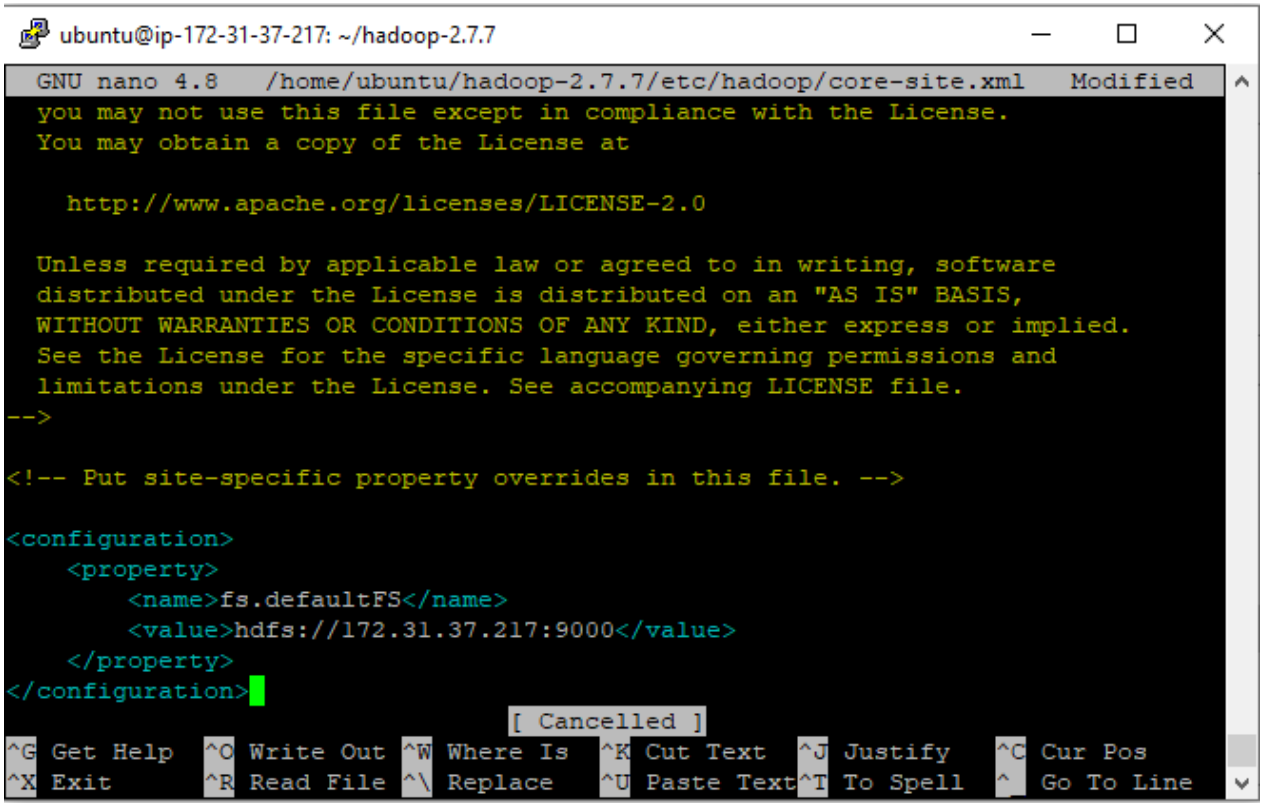

Also used the following command to open the hdfs-site.xml and edited the configuration to adjust the directories of NameNode and DataNode.

- ***sudo nano $HADOOP_HOME/etc/Hadoop/hdfs-site.xml***

```
<configuration>
    <property>
        <name>dfs.replication</name>
        <value>1</value>
    </property>
```

Now format the NameNode before starting Hadoop

- **hdfs namenode -format**

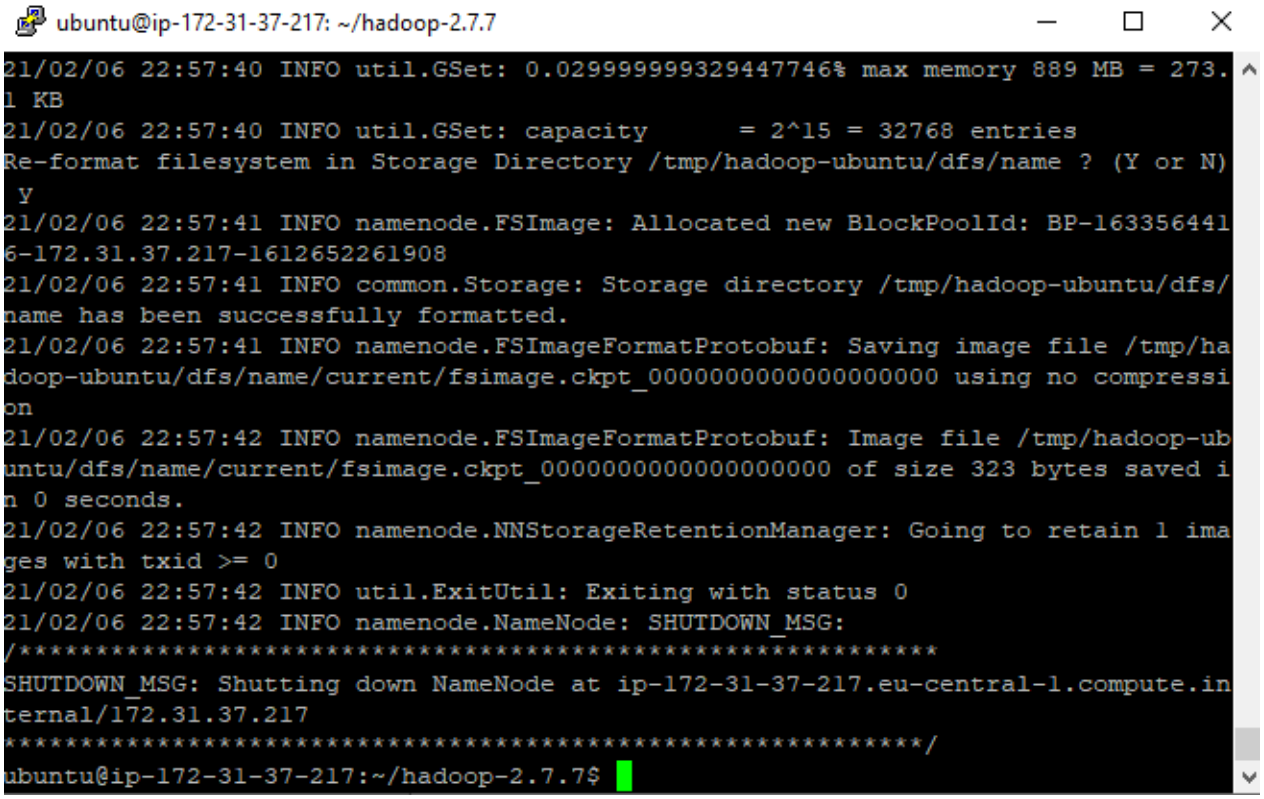

Started DataNode, Yarn resource and namenodes; and checked with command:

- **jps**

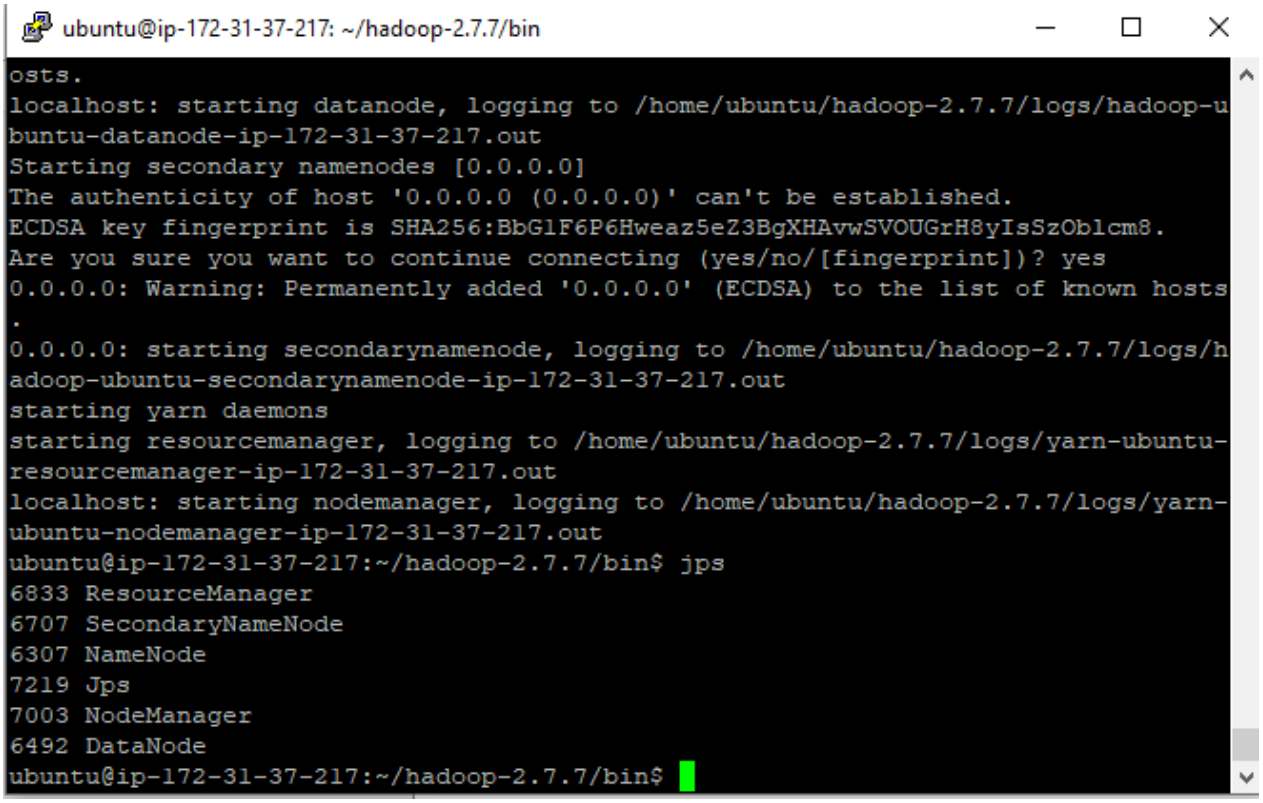

Now our Hadoop environment is ready. Lets move on to the MapReduce WordCounting example.

https://hadoop.apache.org/docs/r2.7.7/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html

This site helped me through this process mostly. I got the *WordCount.java* file from the source code in here. Using **touch WordCount.java** command, I created a java file and with nano command I pasted the source code found in the site.

Project Gutenberg is an online library. So I decided to use it and downloaded the text file of

Frankenstein; Or, The Modern Prometheus by Mary Wollstonecraft Shelley

Using wget with the link of;

- **wget https://www.gutenberg.org/files/84/84-0.txt**

Changed its name to Frankenstein.txt and transferred to Hadoop.

After that I realized that Hadoop and Ubuntu shell language are very similar, it help me a lot to understand the process fast, so using the command below I created the input directory and transferred the text file from my local to HDFS ;

- **hadoop fs –mkdir /WordCountHW**

- **hadoop fs –mkdir /WordCountHW/input**
- **hadoop fs –put Frankenstein.txt /WordCountHW/input**

```
$ bin/hadoop com.sun.tools.javac.Main WordCount.java

$ jar cf wc.jar WordCount*.class
```

Using command below, compiled the java file which contains mapping and reducing elements. And jarred these 3 classes.

- **bin/Hadoop com.sun.tools.jab-vac.Main WordCount.java**
- **jar cf wc.jar WordCount*.class**

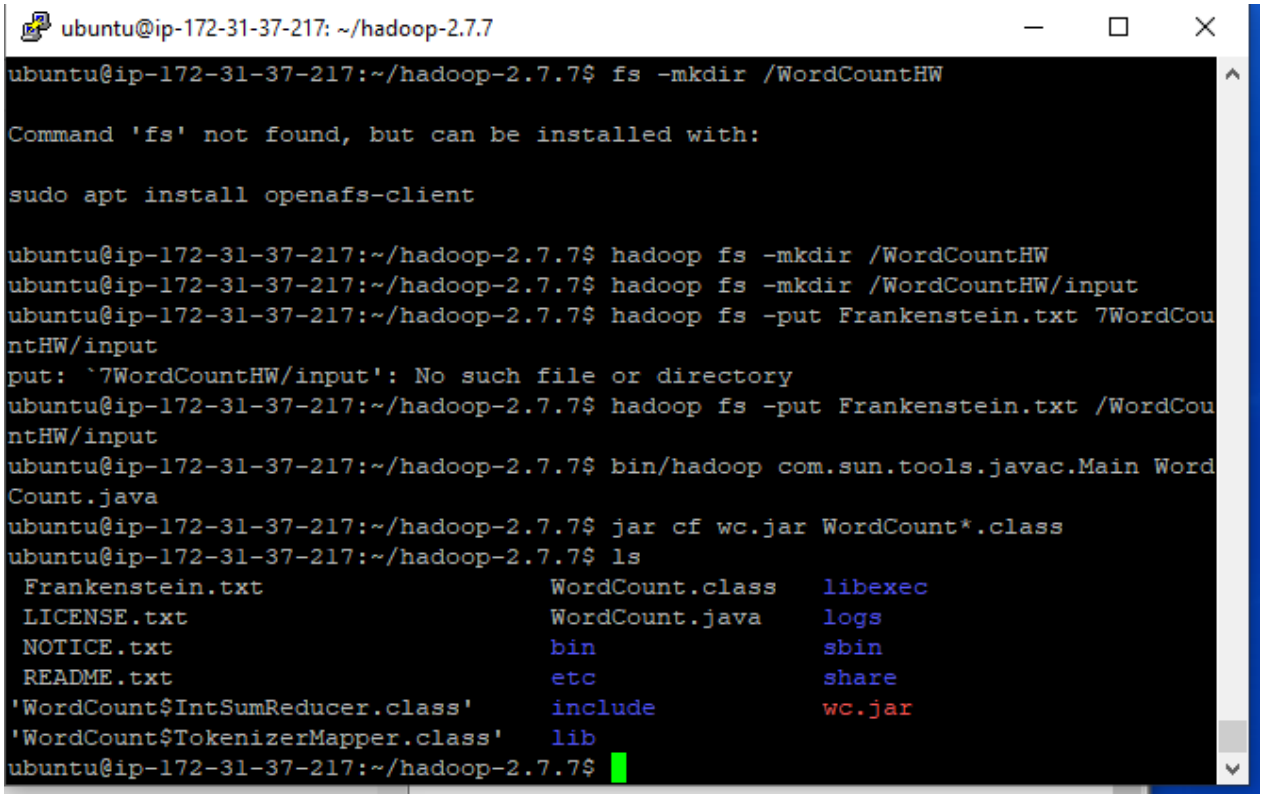

- **bin/hadoop jar wc.jar WordCount /WordCountHW/input /WordCountHW/outputt**

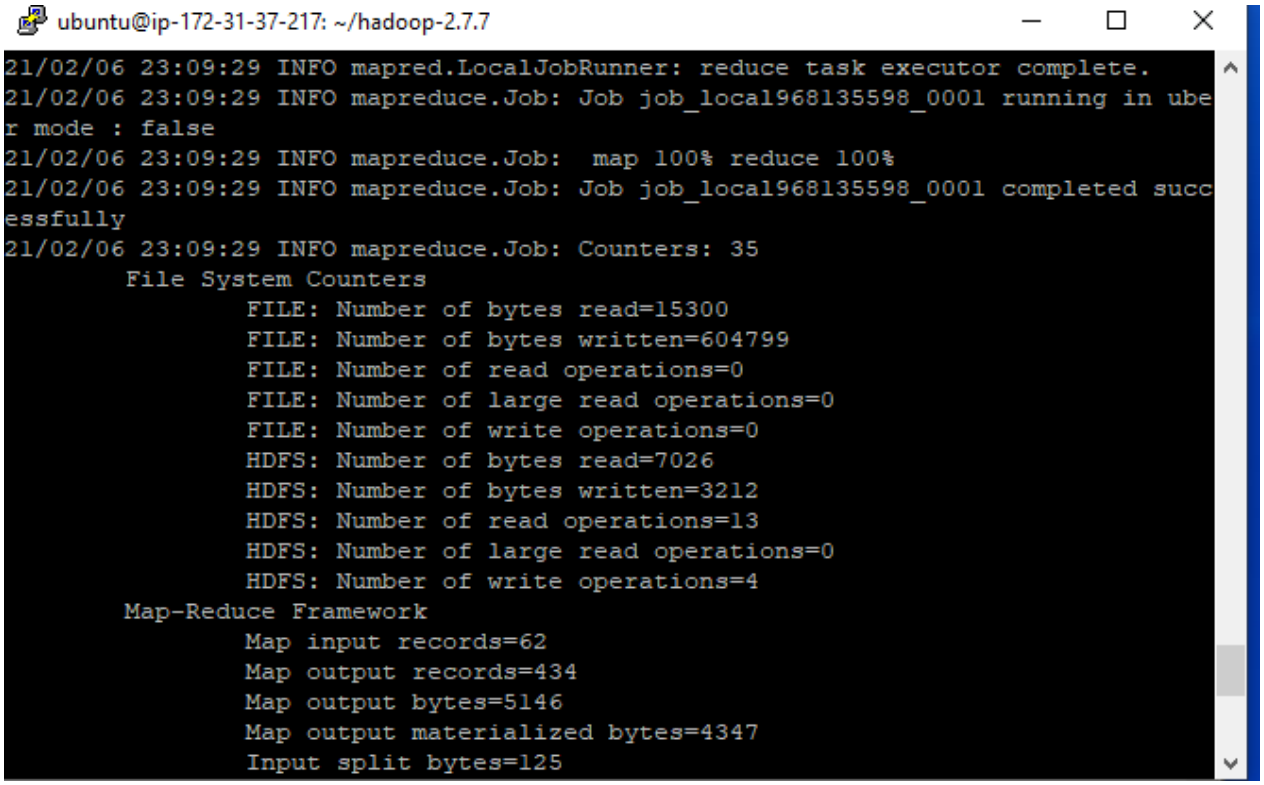

And by the command given below I had the chance to observe the content of the output file;

- **bin/Hadoop fs –cat /WordCountHW/outputt/***

Transferred the file to my local with the command below:

- **hadoop fs –copyToLocal /WordCountHW/outputt/***

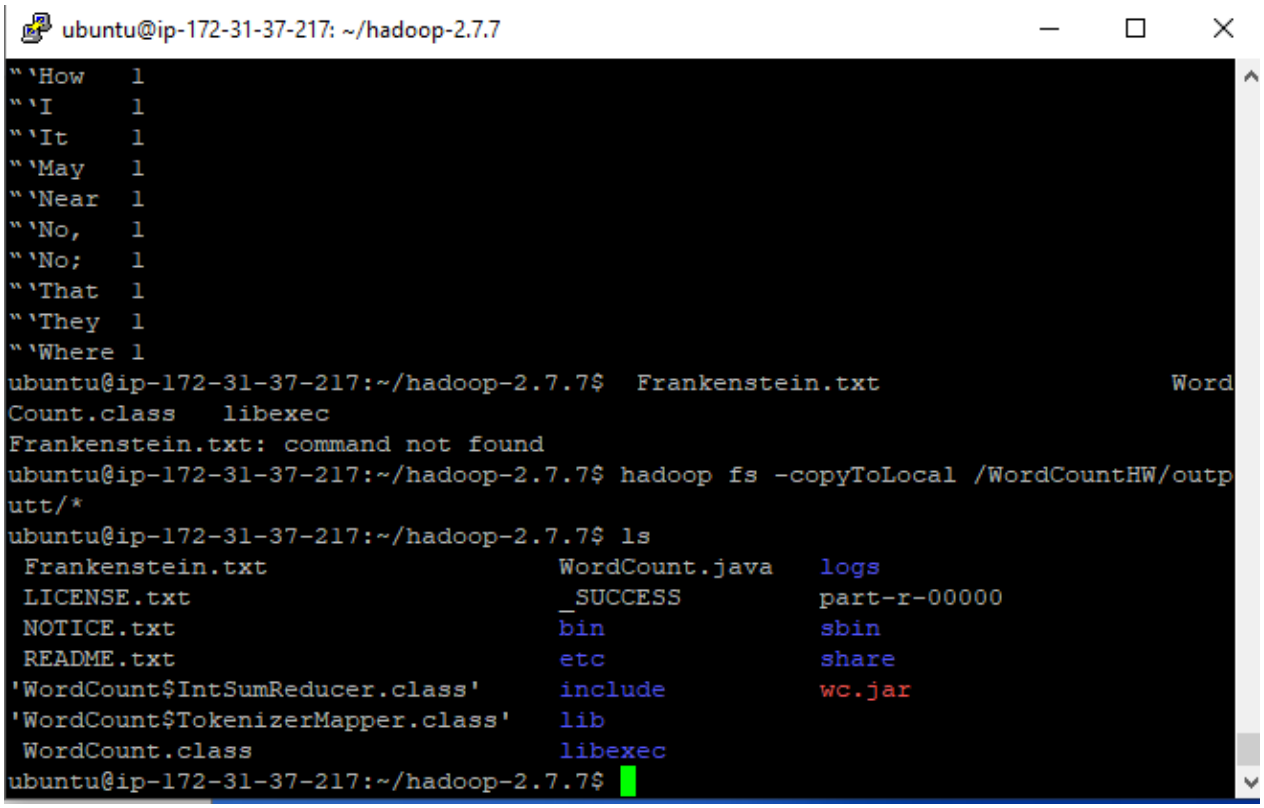

Changed the name of the output file to *FrankensteinWordCount* and using FileZilla Client, transferred the files WordCount.java , classes, Frankenstein input text and output texts to my local. Now sending all of them to you, in case of you want to examine them.